

# 04 - A Typical (Supervised) ML Workflow

ml4econ, HUJI 2023

Itamar Caspi

April 30, 2023 (updated: 2023-05-01)

# Packages and setup

We will use the following packages during the presentation:

```
library(tidyverse)  # for data wrangling and visualization
library(tidymodels) # for data modeling
library(GGally)     # for pairs plot
library(skimr)      # for summary statistics
library(here)       # for referencing folders and files
```

For the presentation, we will select a specific `ggplot` theme (not relevant otherwise):

```
theme_set(theme_grey(20))
```

# The tidymodels package



"**tidymodels** is a "meta-package" for modeling and statistical analysis that share the underlying design philosophy, grammar, and data structures of the tidyverse."

# Supervised Machine Learning Workflow

1. Define the Prediction Task
2. Explore the Data
3. Set Model and Tuning Parameters
4. Perform Cross-Validation
5. Evaluate the Model

# Step 1: Define the Prediction Task

# Welcome to the BostonHousing dataset

- Dataset: 506 census tracts from the 1970 Boston census (Harrison & Rubinfeld, 1978)

Components:

- medv (target): Median home value in thousands of dollars
- lstat (predictor): Percentage of lower status population
- chas (predictor): Proximity to Charles River (1 = yes, 0 = no)

**Objective:** Predict medv based on the given predictors



Source: <https://www.bostonusa.com/>

# A bird's-eye view of Boston



Source: <https://www.wbur.org/news/2019/11/25/heat-mapping-boston-museum-of-science>

# Load the Data

We will utilize the `read_csv()` function to import the raw dataset.

```
boston_raw <- read_csv(here("04-ml-workflow/data", "BostonHousing.csv"))
```



# What Type of Data?

For a better understanding of the data structure, apply the `glimpse()` function:

```
glimpse(boston_raw)
```

```
## Rows: 506
## Columns: 14
## $ crim    <dbl> 0.00632, 0.02731, 0.02729, 0.03237, 0.06905, 0.02985, 0.08829, 0.14455, 0.21124, 0.17~
## $ zn      <dbl> 18.0, 0.0, 0.0, 0.0, 0.0, 0.0, 12.5, 12.5, 12.5, 12.5, 12.5, 12.5, 12.5, 0.0, 0.0, 0.~
## $ indus   <dbl> 2.31, 7.07, 7.07, 2.18, 2.18, 2.18, 7.87, 7.87, 7.87, 7.87, 7.87, 7.87, 7.87, 8.14, 8~
## $ chas    <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0~
## $ nox     <dbl> 0.538, 0.469, 0.469, 0.458, 0.458, 0.458, 0.524, 0.524, 0.524, 0.524, 0.524, 0.524, 0~
## $ rm      <dbl> 6.575, 6.421, 7.185, 6.998, 7.147, 6.430, 6.012, 6.172, 5.631, 6.004, 6.377, 6.009, 5~
## $ age     <dbl> 65.2, 78.9, 61.1, 45.8, 54.2, 58.7, 66.6, 96.1, 100.0, 85.9, 94.3, 82.9, 39.0, 61.8, ~
## $ dis     <dbl> 4.0900, 4.9671, 4.9671, 6.0622, 6.0622, 6.0622, 5.5605, 5.9505, 6.0821, 6.5921, 6.346~
## $ rad     <dbl> 1, 2, 2, 3, 3, 3, 5, 5, 5, 5, 5, 5, 5, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4, 4~
## $ tax     <dbl> 296, 242, 242, 222, 222, 222, 311, 311, 311, 311, 311, 311, 311, 307, 307, 307, 307, ~
## $ ptratio <dbl> 15.3, 17.8, 17.8, 18.7, 18.7, 18.7, 15.2, 15.2, 15.2, 15.2, 15.2, 15.2, 15.2, 21.0, 2~
## $ b       <dbl> 396.90, 396.90, 392.83, 394.63, 396.90, 394.12, 395.60, 396.90, 386.63, 386.71, 392.5~
## $ lstat   <dbl> 4.98, 9.14, 4.03, 2.94, 5.33, 5.21, 12.43, 19.15, 29.93, 17.10, 20.45, 13.27, 15.71, ~
## $ medv    <dbl> 24.0, 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9, 15.0, 18.9, 21.7, 20.4, 1~
```

The `chas` variable predominantly consists of zeros, which implies that it should be treated as a categorical factor.

# Initial Data Filtering

Select medv and lstat

```
boston <- boston_raw %>%  
  as_tibble() %>%  
  select(medv, lstat, chas) %>%  
  mutate(chas = as_factor(chas))  
  
head(boston)
```

```
## # A tibble: 6 x 3  
##   medv lstat chas  
##   <dbl> <dbl> <fct>  
## 1  24    4.98 0  
## 2  21.6   9.14 0  
## 3  34.7   4.03 0  
## 4  33.4   2.94 0  
## 5  36.2   5.33 0  
## 6  28.7   5.21 0
```

## Step 2: Split the Data

# Initial Split

To perform an initial train-test split, we will use the `initial_split()`, `training()`, and `testing()` functions from the **rsample** package.

Remember to set a seed for reproducibility.

```
set.seed(1203)
```

Initial split:

```
boston_split <- boston %>%  
  initial_split(prop = 2/3, strata = medv)
```

```
boston_split
```

```
## <Training/Testing/Total>  
## <336/170/506>
```

# Preparing Training and Test Sets

```
boston_train_raw <- training(boston_split)
boston_test_raw  <- testing(boston_split)

head(boston_train_raw, 5)
```

```
## # A tibble: 5 x 3
##   medv lstat chas
##   <dbl> <dbl> <fct>
## 1  16.5  29.9  0
## 2   15   20.4  0
## 3  13.6  21.0  0
## 4  15.2  18.7  0
## 5  14.5  19.9  0
```

```
head(boston_test_raw, 5)
```

```
## # A tibble: 5 x 3
##   medv lstat chas
##   <dbl> <dbl> <fct>
## 1   24    4.98  0
## 2  21.6    9.14  0
## 3  27.1   19.2  0
## 4  18.9   17.1  0
## 5  18.2   10.3  0
```

# Step 3: Explore the Data

# Summary Statistics Using `skimr`

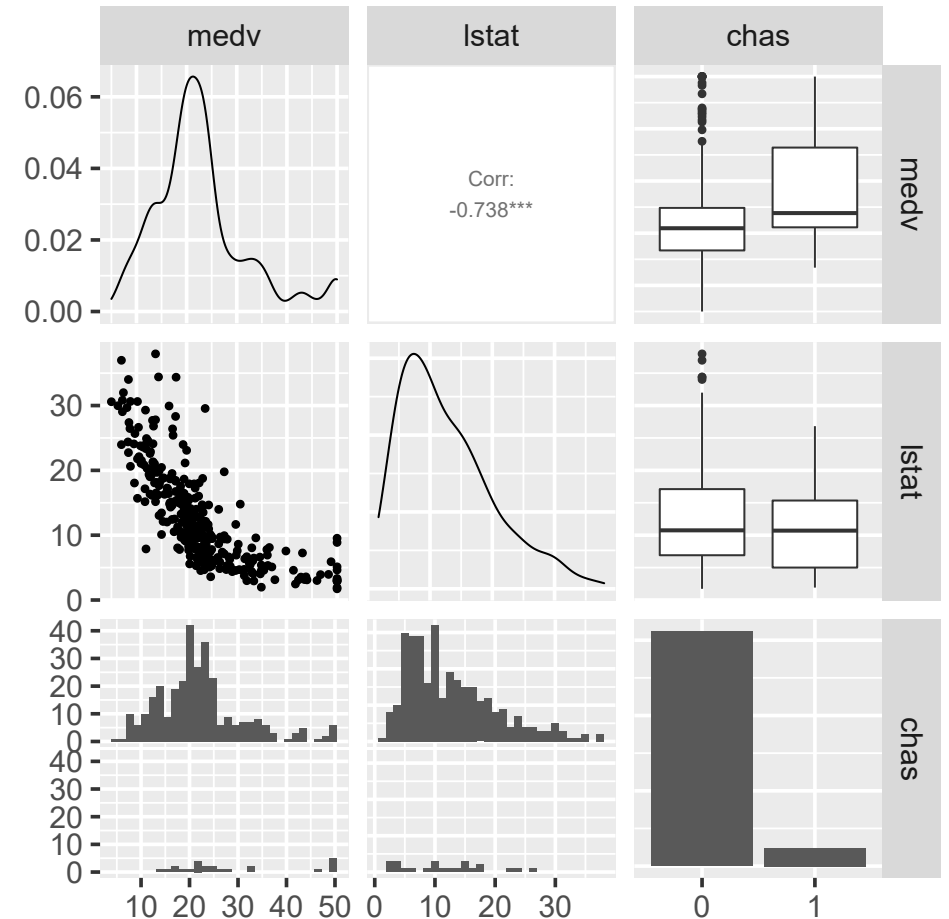
```
boston_train_raw %>%  
  skim()
```

(Not visually appealing on the slides)

# Pairs Plot Using GGally

We will now create a **pairs plot**, which efficiently displays every variable in a dataset against all the others.

```
boston_train_raw %>% ggpairs()
```






# Select a Model

We will select the class of polynomial models, represented as follows:

$$medv_i = \beta_0 + \sum_{j=1}^{\lambda} \beta_j lstat_i^j + \varepsilon_i$$

```
boston_train_raw %>% ggplot(aes(lstat, medv)) +  
  geom_point() +  
  geom_smooth(  
    method = lm,  
    formula = y ~ poly(x,1),  
    se = FALSE,  
    color = "blue"  
  ) +  
  geom_smooth(  
    method = lm,  
    formula = y ~ poly(x,10),  
    se = FALSE,  
    color = "red"  
  )  
)
```

In blue  $\lambda = 1$ ; in red,  $\lambda = 10$ .



## Step 4: Set Model and Tuning Parameters

# Data Preprocessing using recipes

The `recipes` package is an excellent resource for data preprocessing, seamlessly integrating with the tidy approach to machine learning.

```
boston_rec <-  
  recipe(medv ~ lstat + chas, data = boston_train_raw) %>%  
    step_poly(lstat, degree = tune("lambda")) %>%  
    step_dummy(chas)
```

```
boston_rec
```

```
## Recipe  
##  
## Inputs:  
##  
##      role #variables  
## outcome      1  
## predictor      2  
##  
## Operations:  
##  
## Orthogonal polynomials on lstat  
## Dummy variables from chas
```

# Set a Grid for $\lambda$

What are the tuning parameters we need to consider?

```
boston_rec %>% extract_parameter_set_dials()
```

```
## Collection of 1 parameters for tuning  
##  
## identifier  type    object  
##      lambda degree nparam[+]
```

We must tune the polynomial degree parameter ( $\lambda$ ) while constructing our models using the training data. In this example, we will establish a range between 1 and 8:

```
lambda_grid <- expand_grid("lambda" = 1:8)
```

# Define the Model

Using the linear regression model:

```
lm_mod <- linear_reg()%>%  
  set_engine("lm")
```

```
lm_mod
```

```
## Linear Regression Model Specification (regression)  
##  
## Computational engine: lm
```

Note that in this case, there are no tuning parameters involved.

# Step 5: Cross-validation

# Split the Training Set to 5-folds

We will apply the `vfold_cv()` function from the **rsample** package to divide the training set into 5-folds:

```
cv_splits <- boston_train_raw %>%  
  vfold_cv(v = 5)
```

```
cv_splits
```

```
## # 5-fold cross-validation  
## # A tibble: 5 x 2  
##   splits      id  
##   <list>    <chr>  
## 1 <split [268/68]> Fold1  
## 2 <split [269/67]> Fold2  
## 3 <split [269/67]> Fold3  
## 4 <split [269/67]> Fold4  
## 5 <split [269/67]> Fold5
```

# Define the Workflow

Next, we define a `workflow()` that combines a model specification with a recipe or model preprocessor.

```
boston_wf <-  
  workflow() %>%  
  add_model(lm_mod) %>%  
  add_recipe(boston_rec)
```

Note that in this case, there are no tuning parameters involved.



# Estimate CV-RMSE Over the $\lambda$ Grid

We will now calculate the cross-validated root mean squared error (CV-RMSE) for each value of  $\lambda$ .

```
boston_results <-  
  boston_wf %>%  
  tune_grid(  
    resamples = cv_splits,  
    grid      = lambda_grid  
  )  
  
boston_results
```

```
## # Tuning results  
## # 5-fold cross-validation  
## # A tibble: 5 x 4  
##   splits          id   .metrics          .notes  
##   <list>         <chr> <list>          <list>  
## 1 <split [268/68]> Fold1 <tibble [16 x 5]> <tibble [0 x 3]>  
## 2 <split [269/67]> Fold2 <tibble [16 x 5]> <tibble [0 x 3]>  
## 3 <split [269/67]> Fold3 <tibble [16 x 5]> <tibble [0 x 3]>  
## 4 <split [269/67]> Fold4 <tibble [16 x 5]> <tibble [0 x 3]>  
## 5 <split [269/67]> Fold5 <tibble [16 x 5]> <tibble [0 x 3]>
```

# Find the Optimal $\lambda$

Let's identify the top-3 best-performing models.

```
boston_results %>%  
  show_best(metric = "rmse", n = 3)
```


```
## # A tibble: 3 x 7  
##   lambda .metric .estimator mean      n std_err .config  
##   <int> <chr>    <chr>    <dbl> <int>  <dbl> <chr>  
## 1      6 rmse    standard  5.29     5  0.273 Preprocessor6_Model1  
## 2      5 rmse    standard  5.29     5  0.279 Preprocessor5_Model1  
## 3      7 rmse    standard  5.33     5  0.293 Preprocessor7_Model1
```

*"[I]n reality there is rarely if ever a true underlying model, and even if there was a true underlying model, selecting that model will not necessarily give the best forecasts..."*

— **Rob J. Hyndman**

# And Now Using a Graph

```
boston_results %>%  
  autoplot()
```



# Step 6: Evaluate the Model

# Use the Test Set to Evaluate the Best Model

Choose the optimal value of  $\lambda$

```
best_lambda <- boston_results %>%  
  select_best(metric = "rmse")  
  
best_lambda
```

```
## # A tibble: 1 x 2  
##   lambda .config  
##   <int> <chr>  
## 1      6 Preprocessor6_Model1
```

Create a recipe using the optimal  $\lambda = 6$

```
boston_final <- boston_rec %>%  
  finalize_recipe(best_lambda)
```

# Apply the Recipe to the Training and Test Sets

The `juice()` function applies the recipe to the training set, while the `bake()` function applies it to the test set.

```
boston_train <- boston_final %>%  
  prep() %>%  
  juice()  
  
boston_test <- boston_final %>%  
  prep() %>%  
  bake(new_data = boston_test_raw)
```

For instance, let's examine the training set:

```
head(boston_train, 3)
```

```
## # A tibble: 3 x 8  
##   medv lstat_poly_1 lstat_poly_2 lstat_poly_3 lstat_poly_4 lstat_poly_5 lstat_poly_6 chas_X1  
##   <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>      <dbl>  
## 1  16.5         0.126         0.0942      -0.0311      -0.118      -0.0932         0.0108         0  
## 2   15         0.0565        -0.0399      -0.0549       0.0406       0.0604        -0.0342         0  
## 3  13.6         0.0606        -0.0358      -0.0613       0.0335       0.0693        -0.0218         0
```

# Fit the Model to the Training Set

Fit the optimal model (with  $\lambda = 6$ ) to the training set:

```
boston_fit <- lm_mod %>%  
  fit(medv ~ ., data = boston_train)
```

The following are the estimated coefficients:

```
boston_fit %>% tidy()
```

```
## # A tibble: 8 x 5  
##   term          estimate std.error statistic  p.value  
##   <chr>          <dbl>    <dbl>    <dbl>    <dbl>  
## 1 (Intercept)    22.3      0.295     75.6 1.41e-209  
## 2 lstat_poly_1 -126.      5.21    -24.2 4.33e- 75  
## 3 lstat_poly_2   52.8      5.21     10.1 3.49e- 21  
## 4 lstat_poly_3  -21.4      5.23     -4.09 5.36e-  5  
## 5 lstat_poly_4   20.9      5.23      3.99 8.29e-  5  
## 6 lstat_poly_5  -14.7      5.23     -2.80 5.34e-  3  
## 7 lstat_poly_6    4.22      5.22      0.807 4.20e-  1  
## 8 chas_X1        4.45      1.12      3.96 9.27e-  5
```

# Make Predictions Using the Test Set

Generate a tibble that includes the predictions and the actual values:

```
boston_pred <- boston_fit %>%  
  predict(new_data = boston_test) %>%  
  bind_cols(boston_test) %>%  
  select(medv, .pred)  
  
head(boston_pred)
```

```
## # A tibble: 6 x 2  
##   medv .pred  
##   <dbl> <dbl>  
## 1  24    31.3  
## 2  21.6  23.3  
## 3  27.1  15.0  
## 4  18.9  16.7  
## 5  18.2  22.2  
## 6  19.9  24.0
```

It's worth noting that this is the first time we are utilizing the test set!



# Test-RMSE

Calculate the root mean square error (RMSE) for the test set (test-RMSE):

```
boston_pred %>%  
  rmse(medv, .pred)
```

```
## # A tibble: 1 x 3  
##   .metric .estimator .estimate  
##   <chr>   <chr>      <dbl>  
## 1 rmse    standard     5.00
```

The above is a measure of our model's performance on "general" data.


**NOTE:** The test set RMSE estimates the predicted squared error on unseen data, provided the best model.

# Always plot your prediction errors

Plotting the prediction errors ( $y_i - \hat{y}_i$ ) against the target variable provides critical information regarding prediction quality.

```
boston_pred %>%  
  mutate(resid = medv - .pred) %>%  
  ggplot(aes(medv, resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, color = "red")
```

For example, our predictions for high-end levels of `medv` are highly biased, indicating that there's potential for improvement...



# (A shortcut)

The `last_fit()` function from `tune` is a much quicker way to obtain the test-set RMSE.

Firstly, we need to modify our workflow to utilize the optimal  $\lambda$  value.

```
boston_wf <-  
  workflow() %>%  
  add_model(lm_mod) %>%  
  add_recipe(boston_final)
```

We will now use the optimal model to estimate the out-of-sample RMSE.

```
boston_wf %>%  
  last_fit(split = boston_split) %>%  
  collect_metrics() %>%  
  filter(.metric == "rmse")
```

```
## # A tibble: 1 x 4  
##   .metric .estimator .estimate .config  
##   <chr>   <chr>       <dbl> <chr>  
## 1 rmse    standard         5.00 Preprocessor1_Model1
```

```
slides::end()
```

 [Source code](#)